
The Cozo Database Manual

Release 0.1.0

Ziyang Hu

Oct 14, 2022

CONTENTS

1	Introduction	1
2	Getting started	3
2.1	Running Cozo	3
2.2	The query API	3
2.3	Security	4
2.4	Ways of running queries	4
3	Queries	7
3.1	Constant rules	7
3.2	Horn-clause rules	8
3.3	Algorithm application	10
3.4	Query options	10
4	Stored relations and transactions	11
4.1	Using stored relations	11
4.2	Chaining queries into a single transaction	11
4.3	Using transaction to make indices	11
5	Execution of queries	13
6	System ops	15
6.1	Explain	15
6.2	Ops on stored relations	15
6.3	Monitor and kill	15
6.4	Maintenance	15
7	Datatypes	17
7.1	Value types	17
7.2	Value literals	18
7.3	Schema types	18
8	Functions	19
8.1	Equality and Comparisons	19
8.2	Boolean functions	20
8.3	Mathematics	20
8.4	String functions	22
8.5	List functions	23
8.6	Binary functions	24
8.7	Type checking and conversions	25
8.8	Random functions	26

8.9	Regex functions	27
8.10	Timestamp functions	28
9	Aggregations	29
9.1	Meet aggregations	29
9.2	Ordinary aggregations	30
10	Utilities and algorithms	33
10.1	Utilities	33
10.2	Connectedness algorithms	34
10.3	Pathfinding algorithms	35
10.4	Community detection algorithms	37
10.5	Centrality measures	37
10.6	Miscellaneous	39
	Index	41

INTRODUCTION

GETTING STARTED

Cozo is distributed as a single executable. To get started, download the executable for your platform and uncompress it. After decompression, you may also need to give it executable permission by `chmod +x ./cozo` on Unix-based systems.

The pre-compiled distributions of Cozo support Linux, Mac and Windows. As building on Windows is very different from building on UNIX-based systems, the Windows build hasn't received as much attention as the other builds and may suffer from inferior performance and Windows-specific bugs. For Windows users, we recommend running Cozo under [WSL](#)¹ if possible, especially if your workload is heavy.

2.1 Running Cozo

Run the `cozo` command in a terminal:

```
./cozo PATH_TO_DATA_DIRECTORY
```

If `PATH_TO_DATA_DIRECTORY` does not exist, it will be created. Cozo will then start a web server and bind to address `127.0.0.1` and port `9070`. These two can be customized: run the executable with the `-h` option to learn how.

To stop Cozo, type `CTRL-C` in the terminal, or send `SIGTERM` to the process with e.g. `kill`.

2.2 The query API

Queries are run by sending HTTP POST requests to the server. By default, the API endpoint is `http://127.0.0.1:9070/text-query`. The structure of the expected JSON payload is:

```
{
  "script": "<COZOSCRIPT QUERY STRING>",
  "params": {}
}
```

`params` should be an object of named parameters. For example, if you have `params` set up to be `{"num": 1}`, then `$num` can be used anywhere in your query string where an expression is expected. Always use `params` instead of constructing query strings yourself when you have parametrized queries.

¹ <https://learn.microsoft.com/en-us/windows/wsl/install>

2.3 Security

Cozo is currently designed to run in a trusted environment and be used by trusted clients, therefore it does not come with elaborate authentication and security features. If you must access Cozo remotely, you are responsible for setting up firewalls, encryptions and proxies yourself.

As a guard against users carelessly binding Cozo to any address other than `127.0.0.1` and potentially exposing content to everyone on the Internet, in this case, Cozo will refuse to start unless you also set up the environment variable `COZO_AUTH`. With the variable set, Cozo will then require all queries to provide the content of the set variable in the HTTP header field `x-cozo-auth` for verification. Please note that this “security measure” is not considered sufficient for any purpose and is only a last defence when every other security measure that you are responsible for setting up fails.

2.4 Ways of running queries

2.4.1 Making HTTP requests

As Cozo has a web-based API, it is accessible by all languages that are capable of making web requests. The structure of the API is also deliberately kept minimal so that no dedicated clients are necessary. The return values of requests are JSON when requests are successful, or text descriptions when errors occur, so a language only needs to be able to process JSON to use Cozo.

2.4.2 JupyterLab

JupyterLab² is a web-based notebook interface in the python ecosystem heavily used by data scientists and is the recommended “IDE” of Cozo.

First, install JupyterLab by following the install instructions of the project. Then install the `pycozo` library by running:

```
pip install "pycozo[pandas]"
```

Now, open the JupyterLab web interface, start a Python 3 kernel, and in a cell run the following magic command³:

```
%load_ext pycozo.ipyext_direct
```

If you need to connect to Cozo using a non-default address or port, or you require an authentication string, you need to run the following magic commands as well:

```
%cozo_host http://<BIND_ADDRESS>:<PORT>
%cozo_auth <YOUR_AUTH_STRING>
```

Now you can execute cells as you usually do in JupyterLab, and the content of the cells will be sent to Cozo and interpreted as CozoScript. Returned relations will be formatted as `Pandas dataframe`⁴.

The above sets up the notebook in the Direct Cozo mode, where cells are default interpreted as CozoScript. You can still execute python code by starting the first line of a cell with the `%%py`. There is also an Indirect Cozo mode, started by:

² <https://jupyterlab.readthedocs.io/en/stable/>

³ <https://ipython.readthedocs.io/en/stable/interactive/magics.html>

⁴ <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.html>


```
%load_ext pycozo.ipyext
```

In this mode, only cells with the first line content `%%cozo` are interpreted as CozoScript. Other cells are interpreted in the normal way (by default, python code). Which mode you use depends on your workflow. We recommend the Indirect mode if you have lots of post-processing and visualizations.

When a query is successfully executed, the result will be bound to the python variable `_` as a Pandas dataframe (this is a feature of Jupyter notebooks: the Cozo extension didn't do anything extra).

There are a few other useful magic commands:

- `%cozo_run_file <PATH_TO_FILE>` runs a local file as CozoScript.
- `%cozo_run_string <VARIABLE>` runs variable containing string as CozoScript.
- `%cozo_set <KEY> <VALUE>` sets a parameter with the name `<KEY>` to the expression `<VALUE>`. The set parameters will be used by subsequent queries.
- `%cozo_set_params <PARAM_MAP>` replace all parameters by the given expression, which must evaluate to a dictionary with string keys.
- `%cozo_clear` clears all set parameters.
- `%cozo_params` returns the parameters currently set.

2.4.3 The Makeshift JavaScript Console

The Python and JupyterLab ecosystem is rather heavy-weight. If you are just testing out or running Cozo in an environment that only occasionally requires manual queries, you may be reluctant to install them. In this case, you may find the Makeshift JavaScript Console helpful.

As Cozo is running an HTTP service, we assume that the browser on your local machine can reach its network. We recommend [Firefox](https://www.mozilla.org/en-US/firefox/new/)⁵, [Chrome](https://www.google.com/chrome/)⁶, or any Chromium-based browser for best display.

If Cozo is running under the default configuration, navigate to `http://127.0.0.1:9070`. You should be greeted with a mostly empty page telling you that Cozo is running. Now open the Developer Console ([Firefox console](https://firefox-source-docs.mozilla.org/devtools-user/browser_console/index.html)⁷ or [Chrome console](https://developer.chrome.com/docs/devtools/console/javascript/)⁸) and switch to the “Console” tab. Now you can execute CozoScript by running:

```
await run("<COZOSCRIPT>")
```

The returned tables will be properly formatted. If you need to pass in parameters, provide a second parameter with a JavaScript object. If you need to set an auth string, modify the global variable `COZO_AUTH`.

The JavaScript Console is not as nice to use as Jupyter notebooks, but we think that it provides a much better experience than hand-rolled CLI consoles, since you can use JavaScript to manipulate the results.

⁵ <https://www.mozilla.org/en-US/firefox/new/>

⁶ <https://www.google.com/chrome/>

⁷ https://firefox-source-docs.mozilla.org/devtools-user/browser_console/index.html

⁸ <https://developer.chrome.com/docs/devtools/console/javascript/>

QUERIES

The Cozo database system is queried using the CozoScript language. At its core, CozoScript is a [Datalog](#)⁹ dialect supporting stratified negation and stratified recursive meet-aggregations. The built-in native algorithms (mainly graph algorithms) further empower CozoScript for much greater ease of use and much wider applicability.

A query consists of one or many named rules. Each named rule conceptually represents a relation or a table with rows and columns. The rule named ? is called the entry to the query, and its associated relation is returned as the result of the query. Each named rule has associated with it a rule head, which names the columns of the relation, and a rule body, which specifies the content of the relation, or how the content should be computed.

In CozoScript, relations (stored relations or relations defined by rules) abide by the *set semantics*, meaning that even if a rule may compute a row multiple times, it will occur only once in the output. This is in contradistinction to SQL.

There are three types of named rules in CozoScript: constant rules, Horn-clause rules and algorithm applications.

3.1 Constant rules

The following is an example of a constant rule:

```
const_rule[a, b, c] <- [[1, 2, 3], [4, 5, 6]]
```

Constant rules are distinguished by the symbol <- separating the rule head and rule body. The rule body should be an expression evaluating to a list of lists: every sublist of the rule body should be of the same length (the *arity* of the rule), and must match the number of arguments in the rule head. In general, if you are passing data into the query, you should take advantage of named parameters:

```
const_rule[a, b, c] <- $data_passed_in
```

and pass a map containing a key of "data_passed_in" with a value of a list of lists.

The rule head may be omitted if the rule body is not the empty list:

```
const_rule[] <- [[1, 2, 3], [4, 5, 6]]
```

in which case the system will deduce the arity of the rule from the data.

⁹ <https://en.wikipedia.org/wiki/Datalog>

3.2 Horn-clause rules

An example of a Horn-clause rule is:

```
hc_rule[a, e] := rule_a['constant_string', b], rule_b[b, d, a, e]
```

As can be seen, Horn-clause rules are distinguished by the symbol `:=` separating the rule head and rule body. The rule body of a Horn-clause rule consists of multiple *atoms* joined by commas, and is interpreted as representing the *conjunction* of these atoms.

3.2.1 Atoms

Atoms come in various flavours. In the example above:

```
rule_a['constant_string', b]
```

is an atom representing a *rule application*: a rule named `rule_a` must exist in the same query and have the correct arity (2 here). Each row in the named rule is then *unified* with the bindings given as parameters in the square bracket: here the first column is unified with a constant string, and unification succeeds only when the string completely matches what is given; the second column is unified with the *variable* `b`, and as the variable is fresh at this point (meaning that it first appears here), the unification will always succeed and the variable will become *bound*: from this point take on the value of whatever it was unified with in the named relation.

When a bound variable is used again later, for example in `rule_b[b, d, a, e]`, the variable `b` was bound at this point, this unification will only succeed when the unified value is the same as the previously unified value. In other words, repeated use of the same variable in named rules corresponds to inner joins in relational algebra.

Another flavour of atoms is the *stored relation*. It may be written similarly to a rule application:

```
:stored_relation[bind1, bind2]
```

with the colon in front of the stored relation name to distinguish it from rule application. Written in this way, you must give as many bindings to the stored relation as its arity, and the bindings proceed by argument positions, which may be cumbersome and error-prone. So alternatively, you may use the fact that columns of a stored relation are always named and bind by name:

```
:stored_relation{col1: bind1, col2: bind2}
```

In this case, you only need to bind as many variables as you use. If the name you want to give the binding is the same as the name of the column, you may use the shorthand notation: `:stored_relation{col1}` is the same as `:stored_relation{col1: col1}`.

Expressions are also atoms, such as:

```
a > b + 1
```

Here `a` and `b` must be bound somewhere else in the rule, and the expression must evaluate to a boolean, and act as a *filter*: only rows where the expression evaluates to true are kept.

You can also use *unification atoms* to unify explicitly:

```
a = b + c + d
```

for such atoms, whatever appears on the left-hand side must be a single variable and is unified with the right-hand side. This is different from the equality operator `==`, where both sides are merely required to be expressions. When the left-hand side is a single *bound* variable, it may be shown that the equality and the unification operators are semantically equivalent.

Another form of *unification atom* is the explicit multi-unification:

```
a in [x, y, z]
```

here the variable on the left-hand side of `in` is unified with each item on the right-hand side in turn, which in turn implies that the right-hand side must evaluate to a list (but may be represented by a single variable or a function call).

3.2.2 Head and returned relation

Atoms, as explained above, corresponds to either relations (or their projections) or filters in relational algebra. Linked by commas, they, therefore, represent a joined relation, with named columns. The *head* of the rule, which in the simplest case is just a list of variables, then defines whichever columns to keep, and their order in the output relation.

Each variable in the head must be bound in the body, this is one of the *safety rules* of Datalog.

3.2.3 Multiple definitions and disjunction

For Horn-clause rules only, multiple rule definitions may share the same name, with the requirement that the arity of the head in each definition must match. The returned relation is then the *disjunction* of the multiple definitions, which correspond to *union* in SQL. *Intersect* in SQL can be written in CozoScript into a single rule since commas denote conjunction. In complicated situations, you may instead write disjunctions in a single rule with the explicit `or` operator:

```
rule1[a, b] := rule2[a] or rule3[a], rule4[a, b]
```

For completeness, there is also an explicit `and` operator, but it is semantically identical to the comma, except that it has higher operator precedence than `or`, which in turn has higher operator precedence than the comma.

During evaluation, each rule is canonicalized into disjunction normal form and each clause of the outmost disjunction is treated as a separate rule. The consequence is that the safety rule may be violated even though textually every variable in the head occurs in the body. As an example:

```
rule[a, b] := rule1[a] or rule2[b]
```

is a violation of the safety rule since it is rewritten into two rules, each of which is missing a different binding.

3.2.4 Negation

Atoms in Horn clauses may be *negated* by putting `not` in front of them, as in:

```
not rule1[a, b]
```

When negating rule applications and stored relations, at least one binding must be bound somewhere else in the rule in a non-negated context: this is another safety rule of Datalog, and it ensures that the outputs of rules are always finite. The unbound bindings in negated rules remain unbound: negation cannot introduce bound bindings to be used in the head.

Negated expressions act as negative filters, which is semantically equivalent to putting `!` in front of the expression. Since negation does not introduce new bindings, unifications and multi-unifications are converted to equivalent expressions and then negated.

3.2.5 Recursion and stratification

The body of a Horn-clause rule may contain rule applications of itself, and multiple Horn-clause rules may apply each other recursively. The only exception is the entry rule `?`, which cannot be referred to by other rules.

Self and mutual references allow recursion to be defined easily. To guard against semantically pathological cases, recursion cannot occur in negated positions: the Russell-style rule `r[a] := not r[a]` is not allowed. This requirement creates an ordering of the rules, since negated rules must evaluate to completion before rules that apply them can start evaluation: this is called *stratification* of the rules. In cases where a total ordering cannot be defined since there exists a loop in the ordering required by negation, the query is then deemed unstratifiable and Cozo will refuse to execute it.

Note that since CozoScript allows unifying fresh variables, you can still easily write programs that produce infinite relations and hence cannot complete through recursion, but that are still accepted by the database. One of the simplest examples is:

```
r[a] := a = 0
r[a] := r[b], a = b + 1
?[a] := r[a]
```

It is up to the user to ensure that such programs are not submitted to the database, as it is not even in principle possible for the database to rule out such cases without wrongly rejecting valid queries. If you accidentally submitted one, you can refer to the system ops section for how to terminate long-running queries. Or you can give a timeout for the query when you submit.

3.3 Algorithm application

3.4 Query options

STORED RELATIONS AND TRANSACTIONS

4.1 Using stored relations

4.2 Chaining queries into a single transaction

4.3 Using transaction to make indices

EXECUTION OF QUERIES

SYSTEM OPS

6.1 Explain

```
::explain { <query> }
```

6.2 Ops on stored relations

```
::relations  
::relation columns <rel_name>  
::relation remove <rel_name> (, <rel_name>)*  
::relation rename <old_name> -> <new_name> (, <old_name> -> <new_name>)*  
::relation set_triggers <ident> <triggers>  
::relation show_triggers <ident>
```

6.3 Monitor and kill

```
::running  
::kill <pid>
```

6.4 Maintenance

```
::compact
```


DATATYPES

7.1 Value types

A runtime value in Cozo can be of the following *value-types*:

- Null
- Bool
- Number
- String
- Bytes
- Uuid
- List

Number can be Float (double precision) or Int (signed, 64 bits). Cozo will auto-promote Int to Float when necessary.

List can contain any number of mixed-type values, including other lists.

Cozo sorts values according to the above order, e.g. `null` is smaller than `true`, which is in turn smaller than the list `[]`.

Within each type values are *compared* according to logic custom to each type:

- `false < true`;
- `-1 == -1.0 < 0 == 0.0 < 0.5 == 0.5 < 1 == 1.0` (however, see the caveat below);
- Lists are ordered lexicographically by their elements;
- Bytes are compared lexicographically;
- Strings are ordered lexicographically by their UTF-8 byte representations.
- UUIDs are sorted in a way that UUIDv1 with similar timestamps are near each other. This is to improve data locality and should be considered an implementation detail. Depending on the order of UUID in your application is not recommended.

Warning: Because there are two internal number types Int and Float under the umbrella type Number, sorting numbers can be more complex than anticipated.

When sorting, the integer always comes before the equivalent float. For example, `1.0 == 1`, `1.0 >= 1` and `1.0 <= 1` all evaluate to true, but when sorting 1 and 1.0 are two `_different_` values and 1 is placed before 1.0.

This may create problems when applying aggregations since if a grouping key contains both `1.0` and `1`, they are treated as separate group headings. In such cases, it may help to use explicit coercion `to_float` or `round` to coerce all sorted values to the same type.

7.2 Value literals

`null` for the type `Null`, `false` and `true` for the type `Bool` are standard.

Numbers ...

Strings ...

There is no literal representation for `Bytes` due to restrictions placed by JSON. But ...

Lists ...

7.3 Schema types

In schema definition, the required type for a value can be specified by any of the following *schema-types*

- `Ref`
- `Int`
- `Float`
- `Bool`
- `String`
- `Bytes`
- `List`
- `Uuid`

When retrieving triples' values, the schema-type `Ref`, is represented by the type `Uuid`. The entity (the subject of the triple) is always a `Ref`, always represented by a `Uuid`.

Note the absence of the `Null` type in schema-types.

When asserting (inserting or updating) triples, if a value given is not of the correct schema-type, Cozo will first try to coerce the value and will only error out if no known coercion methods exist.

FUNCTIONS

Functions can be used in expressions in Cozo. All function arguments in Cozo are immutable. All functions except those having names starting with `rand_` are deterministic.

Internally, all function arguments are partially evaluated before binding variables to input tuples. For example, the regular expression in `regex_matches(var, '[a-zA-Z]+')` will only be compiled once during the execution of the query, instead of being repeatedly compiled for every input tuple.

8.1 Equality and Comparisons

eq(*x*, *y*)

Equality comparison. The operator form is `x == y` or `x = y`. The two arguments of the equality can be of different types, in which case the result is `false`.

neq(*x*, *y*)

Inequality comparison. The operator form is `x != y`. The two arguments of the equality can be of different types, in which case the result is `true`.

gt(*x*, *y*)

Equivalent to `x > y`

ge(*x*, *y*)

Equivalent to `x >= y`

lt(*x*, *y*)

Equivalent to `x < y`

le(*x*, *y*)

Equivalent to `x <= y`

Note: The four comparison operators can only compare values of the same value type. Integers and floats are of the same type `Number`.

max(*x*, ...)

Returns the maximum of the arguments. Can only be applied to numbers.

min(*x*, ...)

Returns the minimum of the arguments. Can only be applied to numbers.

8.2 Boolean functions

and(...)

Variadic conjunction. For binary arguments it is equivalent to $x \ \&\& \ y$.

or(...)

Variadic disjunction. For binary arguments it is equivalent to $x \ || \ y$.

negate(x)

Negation. Equivalent to $!x$.

assert(x, ...)

Returns true if x is true, otherwise will raise an error containing all its arguments as the error message.

8.3 Mathematics

add(...)

Variadic addition. The binary version is the same as $x + y$.

sub(x, y)

Equivalent to $x - y$.

mul(...)

Variadic multiplication. The binary version is the same as $x * y$.

div(x, y)

Equivalent to x / y .

minus(x)

Equivalent to $-x$.

pow(x, y)

Raises x to the power of y . Equivalent to $x ^ y$. Always returns floating number.

mod(x, y)

Returns the remainder when x is divided by y . Arguments can be floats. The returned value has the same sign as x . Equivalent to $x \% y$.

abs(x)

Returns the absolute value.

signum(x)

Returns 1, 0 or -1, whichever has the same sign as the argument, e.g. `signum(to_float('NEG_INFINITY')) == -1`, `signum(0.0) == 0`, but `signum(-0.0) == -1`. Returns NAN when applied to NAN.

floor(x)

Returns the floor of x .

ceil(x)

Returns the ceiling of x .

round(x)

Returns the nearest integer to the argument (represented as Float if the argument itself is a Float). Round halfway cases away from zero. E.g. `round(0.5) == 1.0`, `round(-0.5) == -1.0`, `round(1.4) == 1.0`.

exp(x)

Returns the exponential of the argument, natural base.

exp2(x)

Returns the exponential base 2 of the argument. Always returns a float.

ln(x)

Returns the natural logarithm.

log2(x)

Returns the logarithm base 2.

log10(x)

Returns the logarithm base 10.

sin(x)

The sine trigonometric Func.

cos(x)

The cosine trigonometric Func.

tan(x)

The tangent trigonometric Func.

asin(x)

The inverse sine.

acos(x)

The inverse cosine.

atan(x)

The inverse tangent.

atan2(x , y)

The inverse tangent [atan2¹⁰](https://en.wikipedia.org/wiki/Atan2) by passing x and y separately.

sinh(x)

The hyperbolic sine.

cosh(x)

The hyperbolic cosine.

tanh(x)

The hyperbolic tangent.

asinh(x)

The inverse hyperbolic sine.

acosh(x)

The inverse hyperbolic cosine.

atanh(x)

The inverse hyperbolic tangent.

deg_to_rad(x)

Converts degrees to radians.

¹⁰ <https://en.wikipedia.org/wiki/Atan2>

rad_to_deg(*x*)

Converts radians to degrees.

haversine(*a_lat*, *a_lon*, *b_lat*, *b_lon*)

Computes with the [haversine formula](#)¹¹ the angle measured in radians between two points *a* and *b* on a sphere specified by their latitudes and longitudes. The inputs are in radians. You probably want the next function since most maps measure angles in radians.

haversine_deg_input(*a_lat*, *a_lon*, *b_lat*, *b_lon*)

Same as the previous function, but the inputs are in degrees instead of radians. The return value is still in radians. If you want the approximate distance measured on the surface of the earth instead of the angle between two points, multiply the result by the radius of the earth, which is about 6371 kilometres, 3959 miles, or 3440 nautical miles.

Warning: The haversine formula, when applied to the surface of the earth, which is not a perfect sphere, can result in an error of less than one percent.

8.4 String functions

length(*str*)

Returns the number of Unicode characters in the string.

Can also be applied to a list or a byte array.

Warning: `length(str)` does not return the number of bytes of the string representation. Also, what is returned depends on the normalization of the string. So if such details are important, apply `unicode_normalize` before `length`.

concat(*x*, ...)

Concatenates strings. Equivalent to `x ++ y` in the binary case.

Can also be applied to lists.

str_includes(*x*, *y*)

Returns `true` if *x* contains the substring *y*, `false` otherwise.

lowercase(*x*)

Convert to lowercase. Supports Unicode.

uppercase(*x*)

Converts to uppercase. Supports Unicode.

trim(*x*)

Removes [whitespace](#)¹² from both ends of the string.

trim_start(*x*)

Removes [whitespace](#)¹³ from the start of the string.

¹¹ https://en.wikipedia.org/wiki/Haversine_formula

¹² https://en.wikipedia.org/wiki/Whitespace_character

¹³ https://en.wikipedia.org/wiki/Whitespace_character

trim_end(*x*)

Removes [whitespace](#)¹⁴ from the end of the string.

starts_with(*x*, *y*)

Tests if *x* starts with *y*.

Tip: `starts_with(var, str)` is preferred over equivalent (e.g. `regex`) conditions, since the compiler may more easily compile the clause into a range scan.

ends_with(*x*, *y*)

tests if *x* ends with *y*.

unicode_normalize(*str*, *norm*)

Converts *str* to the [normalization](#)¹⁵ specified by *norm*. The valid values of *norm* are 'nfc', 'nfd', 'nfkc' and 'nfkd'.

chars(*str*)

Returns Unicode characters of the string as a list of substrings.

from_substrings(*list*)

Combines the strings in *list* into a big string. In a sense, it is the inverse function of `chars`.

Warning: If you want substring slices, indexing strings, etc., first convert the string to a list with `chars`, do the manipulation on the list, and then recombine with `from_substring`. Hopefully, the omission of functions doing such things directly can make people more aware of the complexities involved in manipulating strings (and getting the *correct* result).

8.5 List functions

list(*x*, ...)

Constructs a list from its argument, e.g. `list(1, 2, 3)`. Equivalent to the literal form `[1, 2, 3]`.

is_in(*el*, *list*)

Tests the membership of an element in a list.

first(*l*)

Extracts the first element of the list. Returns `null` if given an empty list.

last(*l*)

Extracts the last element of the list. Returns `null` if given an empty list.

get(*l*, *n*)

Returns the element at index *n* in the list *l*. This function will raise an error if the access is out of bounds. Indices start with 0.

maybe_get(*l*, *n*)

Returns the element at index *n* in the list *l*. This function will return `null` if the access is out of bounds. Indices start with 0.

¹⁴ https://en.wikipedia.org/wiki/Whitespace_character

¹⁵ https://en.wikipedia.org/wiki/Unicode_equivalence

length(*list*)

Returns the length of the list.

Can also be applied to a string or a byte array.

slice(*l, start, end*)

Returns the slice of list between the index *start* (inclusive) and *end* (exclusive). Negative numbers may be used, which is interpreted as counting from the end of the list. E.g. `slice([1, 2, 3, 4], 1, 3) == [2, 3]`, `slice([1, 2, 3, 4], 1, -1) == [2, 3]`.

concat(*x, ...*)

Concatenates lists. The binary case is equivalent to `x ++ y`.

Can also be applied to strings.

prepend(*l, x*)

Prepends *x* to *l*.

append(*l, x*)

Appends *x* to *l*.

reverse(*l*)

Reverses the list.

sorted(*l*)

Sorts the list and returns the sorted copy.

chunks(*l, n*)

Splits the list *l* into chunks of *n*, e.g. `chunks([1, 2, 3, 4, 5], 2) == [[1, 2], [3, 4], [5]]`.

chunks_exact(*l, n*)

Splits the list *l* into chunks of *n*, discarding any trailing elements, e.g. `chunks([1, 2, 3, 4, 5], 2) == [[1, 2], [3, 4]]`.

windows(*l, n*)

Splits the list *l* into overlapping windows of length *n*. e.g. `windows([1, 2, 3, 4, 5], 3) == [[1, 2, 3], [2, 3, 4], [3, 4, 5]]`.

union(*x, y, ...*)

Computes the set-theoretic union of all the list arguments.

intersection(*x, y, ...*)

Computes the set-theoretic intersection of all the list arguments.

difference(*x, y, ...*)

Computes the set-theoretic difference of the first argument with respect to the rest.

8.6 Binary functions

length(*bytes*)

Returns the length of the byte array.

Can also be applied to a list or a string.

bit_and(*x, y*)

Calculate the bitwise and. The two bytes must have the same lengths.

bit_or(*x*, *y*)

Calculate the bitwise or. The two bytes must have the same lengths.

bit_not(*x*)

Calculate the bitwise not.

bit_xor(*x*, *y*)

Calculate the bitwise xor. The two bytes must have the same lengths.

pack_bits([...])packs a list of booleans into a byte array; if the list is not divisible by 8, it is padded with `false`.**unpack_bits**(*x*)

Unpacks a byte array into a list of booleans.

encode_base64(*b*)Encodes the byte array *b* into the [Base64](#)¹⁶-encoded string.

Note: `encode_base64` is automatically applied when output to JSON since JSON cannot represent bytes natively.

decode_base64(*str*)Tries to decode the *str* as a [Base64](#)¹⁷-encoded byte array.

8.7 Type checking and conversions

to_string(*x*)Convert *x* to a string: the argument is unchanged if it is already a string, otherwise its JSON string representation will be returned.**to_float**(*x*)Tries to convert *x* to a float. Conversion from numbers always succeeds. Conversion from strings has the following special cases in addition to the usual string representation:

- `INF` is converted to infinity;
- `NEG_INF` is converted to negative infinity;
- `NAN` is converted to `NAN` (but don't compare `NAN` by equality, use `is_nan` instead);
- `PI` is converted to `pi` (3.14159...);
- `E` is converted to the base of natural logarithms, or Euler's constant (2.71828...).

to_uuid(*x*)Tries to convert *x* to a UUID. The input must either be a hyphenated UUID string representation or already a UUID for it to succeed.**uuid_timestamp**(*x*)Extracts the timestamp from a UUID version 1, as seconds since the UNIX epoch. If the UUID is not of version 1, `null` is returned. If *x* is not a UUID, an error is raised.**is_null**(*x*)Checks for `null`.

¹⁶ <https://en.wikipedia.org/wiki/Base64>¹⁷ <https://en.wikipedia.org/wiki/Base64>

is_int(*x*)

Checks for integers.

is_float(*x*)

Checks for floats.

is_finite(*x*)

Returns **true** if *x* is an integer or a finite float.

is_infinite(*x*)

Returns **true** if *x* is infinity or negative infinity.

is_nan(*x*)

Returns **true** if *x* is the special float **NAN**. Returns **false** when the argument is not of number type.

is_num(*x*)

Checks for numbers.

is_bytes(*x*)

Checks for bytes.

is_list(*x*)

Checks for lists.

is_string(*x*)

Checks for strings.

is_uuid(*x*)

Checks for UUIDs.

8.8 Random functions

rand_float()

Generates a float in the interval [0, 1], sampled uniformly.

rand_bernoulli(*p*)

Generates a boolean with probability *p* of being **true**.

rand_int(*lower*, *upper*)

Generates an integer within the given bounds, both bounds are inclusive.

rand_choose(*list*)

Randomly chooses an element from *list* and returns it. If the list is empty, it returns **null**.

rand_uuid_v1()

Generate a random UUID, version 1 (random bits plus timestamp).

rand_uuid_v4()

Generate a random UUID, version 4 (completely random bits).

8.9 Regex functions

regex_matches(*x*, *reg*)

Tests if *x* matches the regular expression *reg*.

regex_replace(*x*, *reg*, *y*)

Replaces the first occurrence of the pattern *reg* in *x* with *y*.

regex_replace_all(*x*, *reg*, *y*)

Replaces all occurrences of the pattern *reg* in *x* with *y*.

regex_extract(*x*, *reg*)

Extracts all occurrences of the pattern *reg* in *x* and returns them in a list.

regex_extract_first(*x*, *reg*)

Extracts the first occurrence of the pattern *reg* in *x* and returns it. If none is found, returns null.

8.9.1 Regex syntax

Matching one character:

.	any character except new line
\d	digit (\p{Nd})
\D	not digit
\pN	One-letter name Unicode character class
\p{Greek}	Unicode character class (general category or script)
\PN	Negated one-letter name Unicode character class
\P{Greek}	negated Unicode character class (general category or script)

Character classes:

[xyz]	A character class matching either x, y or z (union).
[^xyz]	A character class matching any character except x, y and z.
[a-z]	A character class matching any character in range a-z.
[[:alpha:]]	ASCII character class ([A-Za-z])
[[:^alpha:]]	Negated ASCII character class ([^A-Za-z])
[x[^xyz]]	Nested/grouping character class (matching any character except y and z)
[a-y&&xyz]	Intersection (matching x or y)
[0-9&&[^4]]	Subtraction using intersection and negation (matching 0-9 except 4)
[0-9--4]	Direct subtraction (matching 0-9 except 4)
[a-g~b-h]	Symmetric difference (matching `a` and `h` only)
[\[\]]	Escaping in character classes (matching [or])

Composites:

xy	concatenation (x followed by y)
x y	alternation (x or y, prefer x)

Repetitions:

x*	zero or more of x (greedy)
x+	one or more of x (greedy)
x?	zero or one of x (greedy)

(continues on next page)

(continued from previous page)

<code>x*?</code>	zero or more of <code>x</code> (ungreedy/lazy)
<code>x+?</code>	one or more of <code>x</code> (ungreedy/lazy)
<code>x??</code>	zero or one of <code>x</code> (ungreedy/lazy)
<code>x{n,m}</code>	at least <code>n</code> <code>x</code> and at most <code>m</code> <code>x</code> (greedy)
<code>x{n,}</code>	at least <code>n</code> <code>x</code> (greedy)
<code>x{n}</code>	exactly <code>n</code> <code>x</code>
<code>x{n,m}?</code>	at least <code>n</code> <code>x</code> and at most <code>m</code> <code>x</code> (ungreedy/lazy)
<code>x{n,}?</code>	at least <code>n</code> <code>x</code> (ungreedy/lazy)
<code>x{n}?</code>	exactly <code>n</code> <code>x</code>

Empty matches:

<code>^</code>	the beginning of the text
<code>\$</code>	the end of the text
<code>\A</code>	only the beginning of the text
<code>\Z</code>	only the end of the text
<code>\b</code>	a Unicode word boundary (<code>\w</code> on one side and <code>\W</code> , <code>\A</code> , or <code>\Z</code> on the other)
<code>\B</code>	not a Unicode word boundary

8.10 Timestamp functions

now()

Returns the current timestamp as seconds since the UNIX epoch.

format_timestamp(*ts*, *tz*?)

Interpret *ts* as seconds since the epoch and format as a string according to RFC3339¹⁸.

If a second string argument is provided, it is interpreted as a [timezone](#)¹⁹ and used to format the timestamp.

parse_timestamp(*str*)

Parse *str* into seconds since the epoch according to RFC3339.

¹⁸ <https://www.rfc-editor.org/rfc/rfc3339>

¹⁹ https://en.wikipedia.org/wiki/Tz_database

AGGREGATIONS

Aggregations in Cozo can be thought of as a function that acts on a string of values and produces a single value (the aggregate). Due to Datalog semantics, the stream is never empty.

There are two kinds of aggregations in Cozo, *ordinary aggregations* and *meet aggregations*. They are implemented differently in Cozo, with meet aggregations generally faster and more powerful (e.g. only meet aggregations can be recursive).

The power of meet aggregations derive from the additional properties they satisfy by forming a [semilattice](https://en.wikipedia.org/wiki/Semilattice)²⁰:

idempotency

the aggregate of a single value *a* is *a* itself,

commutivity

the aggregate of *a* then *b* is equal to the aggregate of *b* then *a*,

associativity

it is immaterial where we put the parentheses in an aggregate application.

Meet aggregations can be used as ordinary ones, but the reverse is impossible.

9.1 Meet aggregations

min(*x*)

Aggregate the minimum value of all *x*.

max(*x*)

Aggregate the maximum value of all *x*.

and(*var*)

Aggregate the logical conjunction of the variable passed in.

or(*var*)

Aggregate the logical disjunction of the variable passed in.

union(*var*)

Aggregate the unions of *var*, which must be a list.

intersection(*var*)

Aggregate the intersections of *var*, which must be a list.

²⁰ <https://en.wikipedia.org/wiki/Semilattice>

choice(*var*)

Non-deterministically chooses one of the values of *var* as the aggregate. It simply chooses the first value it meets (the order that it meets values should be considered non-deterministic).

choice_last(*var*)

Non-deterministically chooses one of the values of *var* as the aggregate. It simply chooses the last value it meets.

min_cost([*data*, *cost*])

The argument should be a list of two elements and this aggregation chooses the list of the minimum *cost*.

shortest(*var*)

var must be a list. Returns the shortest list among all values. Ties will be broken non-deterministically.

coalesce(*var*)

Returns the first non-null value it meets. The order is non-deterministic.

bit_and(*var*)

var must be bytes. Returns the bitwise ‘and’ of the values.

bit_or(*var*)

var must be bytes. Returns the bitwise ‘or’ of the values.

9.2 Ordinary aggregations

count(*var*)

Count how many values are generated for *var* (using bag instead of set semantics).

count_unique(*var*)

Count how many unique values there are for *var*.

collect(*var*)

Collect all values for *var* into a list.

unique(*var*)

Collect *var* into a list, keeping each unique value only once.

group_count(*var*)

Count the occurrence of unique values of *var*, putting the result into a list of lists, e.g. when applied to 'a', 'b', 'c', 'c', 'a', 'c', the results is [['a', 2], ['b', 1], ['c', 3]].

bit_xor(*var*)

var must be bytes. Returns the bitwise ‘xor’ of the values.

latest_by([*data*, *time*])

The argument should be a list of two elements and this aggregation returns the data of the maximum cost. This is very similar to *min_cost*, the differences being that maximum instead of minimum is used, only the data itself is returned, and the aggregation is deliberately not a meet aggregation. Intended to be used in timestamped audit trails.

9.2.1 Statistical aggregations

mean(x)

The mean value of x .

sum(x)

The sum of x .

product(x)

The product of x .

variance(x)

The sample variance of x .

std_dev(x)

The sample standard deviation of x .

UTILITIES AND ALGORITHMS

The purpose of the native, built-in algorithms in Cozo is to enable easy computation of results that would require either queries too awkward to express in pure Datalog, or time or space requirements that are unreasonable if implemented in the interpreted queries framework. More algorithms will be added as Cozo is further developed.

10.1 Utilities

ReorderSort(*rel*[...], *out*: [...], *sort_by*: [...], *descending*: *false*, *break_ties*: *false*, *skip*: 0, *take*: 0)

Sort and then extract new columns of the passed in relation *rel*.

Parameters

- **out** (*required*) – A list of expressions which will be used to produce the output relation. Any bindings in the expressions will be bound to the named positions in *rel*.
- **sort_by** – A list of expressions which will be used to produce the sort keys. Any bindings in the expressions will be bound to the named positions in *rel*.
- **descending** – Whether the sorting process should be done in descending order. Defaults to *false*.
- **break_ties** – Whether ties should be broken, e.g. whether the first two rows with *identical* sort keys should be given ordering numbers 1 and 2 instead of 1 and 1. Defaults to *false*.
- **skip** – How many rows to skip before producing rows. Defaults to zero.
- **take** – How many rows at most to produce. Zero means no limit. Defaults to zero.

Returns

The returned relation, in addition to the rows specified in the parameter *out*, will have the ordering prepended. The ordering starts at 1.

Tip: This algorithm serves a similar purpose to the global `:order`, `:limit` and `:offset` options, but can be applied to intermediate results. Prefer the global options if it is applied to the final output.

CsvReader(*url*: ..., *types*: [...], *delimiter*: ',', *prepend_index*: *false*, *has_headers*: *true*)

Read a CSV file from disk or an HTTP GET request and convert the result to a relation.

Parameters

- **url** (*required*) – URL for the CSV file. For local file, use `file://<PATH_TO_FILE>`.
- **types** (*required*) – A list of strings interpreted as types for the columns of the output relation. If any type is specified as nullable and conversion to the specified type fails, `null`

will be the result. This is more lenient than other functions since CSVs tend to contain lots of bad values.

- **delimiter** – The delimiter to use when parsing the CSV file.
- **prepend_index** – If `true`, row index will be prepended to the columns.
- **has_headers** – Whether the CSV file has headers. The reader will not interpret the header in any way but will instead simply ignore it.

JsonReader(*url*: ..., *fields*: [...], *json_lines*: *true*, *null_if_absent*: *false*, *prepend_index*: *false*)

Read a JSON file for disk or an HTTP GET request and convert the result to a relation.

Parameters

- **url** (*required*) – URL for the JSON file. For local file, use `file://<PATH_TO_FILE>`.
- **fields** (*required*) – A list of field names, for extracting fields from JSON arrays into the relation.
- **json_lines** – If `true`, parse the file as lines of JSON objects, each line containing a single object; if `false`, parse the file as a JSON array containing many objects.
- **null_if_absent** – If a `true` and a requested field is absent, will output `null` in its place. If `false` and the requested field is absent, will throw an error.
- **prepend_index** – If `true`, row index will be prepended to the columns.

10.2 Connectedness algorithms

ConnectedComponents(*edges*[*from*, *to*])

Computes the [connected components](#)²¹ of a graph with the provided edges.

Returns

Pairs containing the node index, and its component index.

StronglyConnectedComponent(*edges*[*from*, *to*])

Computes the [strongly connected components](#)²² of a graph with the provided edges.

Returns

Pairs containing the node index, and its component index.

SCC(...)

See *Algo.StronglyConnectedComponent*.

MinimumSpanningForestKruskal(*edges*[*from*, *to*, *weight*?])

Runs [Kruskal's algorithm](#)²³ on the provided edges to compute a [minimum spanning forest](#)²⁴. Negative weights are fine.

Returns

Triples containing the from-node, the to-node, and the cost from the tree root to the to-node. Which nodes are chosen to be the roots are non-deterministic. Multiple roots imply the graph is disconnected.

²¹ [https://en.wikipedia.org/wiki/Connected_component_\(graph_theory\)](https://en.wikipedia.org/wiki/Connected_component_(graph_theory))

²² https://en.wikipedia.org/wiki/Strongly_connected_component

²³ https://en.wikipedia.org/wiki/Kruskal%27s_algorithm

²⁴ https://en.wikipedia.org/wiki/Minimum_spanning_tree

MinimumSpanningTreePrim(*edges*[*from*, *to*, *weight*?], *starting*?[*idx*])

Runs [Prim's algorithm](#)²⁵ on the provided edges to compute a [minimum spanning tree](#)²⁶. *starting* should be a relation producing exactly one node index as the starting node. Only the connected component of the starting node is returned. If *starting* is omitted, which component is returned is arbitrary.

Returns

Triples containing the from-node, the to-node, and the cost from the tree root to the to-node.

TopSort(*edges*[*from*, *to*])

Performs [topological sorting](#)²⁷ on the graph with the provided edges. The graph is required to be connected in the first place.

Returns

Pairs containing the sort order and the node index.

10.3 Pathfinding algorithms

ShortestPathDijkstra(*edges*[*from*, *to*, *weight*?], *starting*[*idx*], *goals*[*idx*], *undirected*: *false*, *keep_ties*: *false*)

Runs [Dijkstra's algorithm](#)²⁸ to determine the shortest paths between the *starting* nodes and the *goals*. Weights, if given, must be non-negative.

Parameters

- **undirected** – Whether the graph should be interpreted as undirected. Defaults to *false*.
- **keep_ties** – Whether to return all paths with the same lowest cost. Defaults to *false*, in which any one path of the lowest cost could be returned.

Returns

4-tuples containing the starting node, the goal, the lowest cost, and a path with the lowest cost.

KShortestPathYen(*edges*[*from*, *to*, *weight*?], *starting*[*idx*], *goals*[*idx*], *k*: *expr*, *undirected*: *false*)

Runs [Yen's algorithm](#)²⁹ (backed by Dijkstra's algorithm) to find the k-shortest paths between nodes in *starting* and nodes in *goals*.

Parameters

- **k** (*required*) – How many routes to return for each start-goal pair.
- **undirected** – Whether the graph should be interpreted as undirected. Defaults to *false*.

Returns

4-tuples containing the starting node, the goal, the cost, and a path with the cost.

BreadthFirstSearch(*edges*[*from*, *to*], *nodes*[*idx*, ...], *starting*?[*idx*], *condition*: *expr*, *limit*: *1*)

Runs breadth first search on the directed graph with the given edges and nodes, starting at the nodes in *starting*. If *starting* is not given, it will default to all of nodes, which may be quite a lot to calculate.

Parameters

- **condition** (*required*) – The stopping condition, will be evaluated with the bindings given to nodes. Should evaluate to a boolean, with *true* indicating an acceptable answer was found.

²⁵ https://en.wikipedia.org/wiki/Prim%27s_algorithm

²⁶ https://en.wikipedia.org/wiki/Minimum_spanning_tree

²⁷ https://en.wikipedia.org/wiki/Topological_sorting

²⁸ https://en.wikipedia.org/wiki/Dijkstra%27s_algorithm

²⁹ https://en.wikipedia.org/wiki/Yen%27s_algorithm

- **limit** – How many answers to produce for each starting nodes. Defaults to 1.

Returns

Triples containing the starting node, the answer node, and the found path connecting them.

BFS(...)

See [Algo.BreadthFirstSearch](#).

DepthFirstSearch(*edges*[*from*, *to*], *nodes*[*idx*, ...], *starting*?[*idx*], *condition*: *expr*, *limit*: 1)

Runs depth first search on the directed graph with the given edges and nodes, starting at the nodes in **starting**. If **starting** is not given, it will default to all of **nodes**, which may be quite a lot to calculate.

Parameters

- **condition** (*required*) – The stopping condition, will be evaluated with the bindings given to **nodes**. Should evaluate to a boolean, with **true** indicating an acceptable answer was found.
- **limit** – How many answers to produce for each starting nodes. Defaults to 1.

Returns

Triples containing the starting node, the answer node, and the found path connecting them.

Tip: You probably don't want to use depth first search for path finding unless you have a really niche use case.

DFS(...)

See [Algo.DepthFirstSearch](#).

ShortestPathAStar(*edges*[*from*, *to*, *weight*], *nodes*[*idx*, ...], *starting*[*idx*], *goals*[*idx*], *heuristic*: *expr*)

Computes the shortest path from every node in **starting** to every node in **goals** by the A* algorithm³⁰. **edges** are interpreted as directed, weighted edges with non-negative weights.

Parameters

heuristic (*required*) – The search heuristic expression. It will be evaluated with the bindings from **goals** and **nodes**. It should return a number which is a lower bound of the true shortest distance from a node to the goal node. If the estimate is not a valid lower-bound, i.e. it over-estimates, the results returned may not be correct.

Returns

4-tuples containing the starting node index, the goal node index, the lowest cost, and a path with the lowest cost.

Tip: The performance of A* star algorithm heavily depends on how good your heuristic function is. Passing in 0 as the estimate is always valid, but then you really should be using Dijkstra's algorithm.

Good heuristics usually come about from a metric in the ambient space in which your data live, e.g. spherical distance on the surface of a sphere, or Manhattan distance on a grid. [Func.Math.haversine_deg_input](#) could be helpful for the spherical case. Note that you must use the correct units for the distance.

Providing a heuristic that is not guaranteed to be a lower-bound *might* be acceptable if you are fine with inaccuracies. The errors in the answers are bound by the sum of the margins of your over-estimates.

³⁰ https://en.wikipedia.org/wiki/A*_search_algorithm

10.4 Community detection algorithms

ClusteringCoefficients(*edges[from, to, weight?]*)

Computes the [clustering coefficients](#)³¹ of the graph with the provided edges.

Returns

4-tuples containing the node index, the clustering coefficient, the number of triangles attached to the node, and the total degree of the node.

CommunityDetectionLouvain(*edges[from, to, weight?], undirected: false, max_iter: 10, delta: 0.0001, keep_depth?: depth*)

Runs the [Louvain algorithm](#)³² on the graph with the provided edges, optionally non-negatively weighted.

Parameters

- **undirected** – Whether the graph should be interpreted as undirected. Defaults to `false`.
- **max_iter** – The maximum number of iterations to run within each epoch of the algorithm. Defaults to 10.
- **delta** – How much the [modularity](#)³³ has to change before a step in the algorithm is considered to be an improvement.
- **keep_depth** – How many levels in the hierarchy of communities to keep in the final result. If omitted, all levels are kept.

Returns

Pairs containing the label for a community, and a node index belonging to the community. Each label is a list of integers with maximum length constrained by the parameter `keep_depth`. This list represents the hierarchy of sub-communities containing the list.

LabelPropagation(*edges[from, to, weight?], undirected: false, max_iter: 10*)

Runs the [label propagation algorithm](#)³⁴ on the graph with the provided edges, optionally weighted.

Parameters

- **undirected** – Whether the graph should be interpreted as undirected. Defaults to `false`.
- **max_iter** – The maximum number of iterations to run. Defaults to 10.

Returns

Pairs containing the integer label for a community, and a node index belonging to the community.

10.5 Centrality measures

DegreeCentrality(*edges[from, to]*)

Computes the degree centrality of the nodes in the graph with the given edges. The computation is trivial, so this should be your first thing to try when exploring new data.

Returns

4-tuples containing the node index, the total degree (how many edges involve this node), the out-degree (how many edges point away from this node), and the in-degree (how many edges point to this node).

³¹ https://en.wikipedia.org/wiki/Clustering_coefficient

³² https://en.wikipedia.org/wiki/Louvain_method

³³ [https://en.wikipedia.org/wiki/Modularity_\(networks\)](https://en.wikipedia.org/wiki/Modularity_(networks))

³⁴ https://en.wikipedia.org/wiki/Label_propagation_algorithm

PageRank(*edges[from, to, weight?], undirected: false, theta: 0.8, epsilon: 0.05, iterations: 20*)

Computes the [PageRank](#)³⁵ from the given graph with the provided edges, optionally weighted.

Parameters

- **undirected** – Whether the graph should be interpreted as undirected. Defaults to *false*.
- **theta** – A number between 0 and 1 indicating how much weight in the PageRank matrix is due to the explicit edges. A number of 1 indicates no random restarts. Defaults to 0.8.
- **epsilon** – Minimum PageRank change in any node for an iteration to be considered an improvement. Defaults to 0.05.
- **iterations** – How many iterations to run. Fewer iterations are run if convergence is reached. Defaults to 20.

Returns

Pairs containing the node label and its PageRank. For a graph with uniform edges, the PageRank of every node is 1. The [L2-norm](#)³⁶ of the results is forced to be invariant, i.e. in the results those nodes with a PageRank greater than 1 is “more central” than the average node in a certain sense.

ClosenessCentrality(*edges[from, to, weight?], undirected: false*)

Computes the [closeness centrality](#)³⁷ of the graph. The input relation represent edges connecting node indices which are optionally weighted.

Parameters

undirected – Whether the edges should be interpreted as undirected. Defaults to *false*.

Returns

Node index together with its centrality.

BetweennessCentrality(*edges[from, to, weight?], undirected: false*)

Computes the [betweenness centrality](#)³⁸ of the graph. The input relation represent edges connecting node indices which are optionally weighted.

Parameters

undirected – Whether the edges should be interpreted as undirected. Defaults to *false*.

Returns

Node index together with its centrality.

Warning: `BetweennessCentrality` is very expensive to compute for medium to large graphs. Plan resources accordingly.

³⁵ <https://en.wikipedia.org/wiki/PageRank>

³⁶ [https://en.wikipedia.org/wiki/Norm_\(mathematics\)](https://en.wikipedia.org/wiki/Norm_(mathematics))

³⁷ https://en.wikipedia.org/wiki/Closeness_centrality

³⁸ https://en.wikipedia.org/wiki/Betweenness_centrality

10.6 Miscellaneous

RandomWalk(*edges*[*from*, *to*, ...], *nodes*[*idx*, ...], *starting*[*idx*], *steps*: 10, *weight*?: *expr*, *iterations*: 1)

Performs random walk on the graph with the provided edges and nodes, starting at the nodes in *starting*.

Parameters

- **steps** (*required*) – How many steps to walk for each node in *starting*. Produced paths may be shorter if dead ends are reached.
- **weight** – An expression evaluated against bindings of *nodes* and bindings of *edges*, at a time when the walk is at a node and choosing between multiple edges to follow. It should evaluate to a non-negative number indicating the weight of the given choice of edge to follow. If omitted, which edge to follow is chosen uniformly.
- **iterations** – How many times walking is repeated for each starting node.

Returns

Triples containing a numerical index for the walk, the starting node, and the path followed.

A

abs() (in module *Func.Math*), 20
acos() (in module *Func.Math*), 21
acosh() (in module *Func.Math*), 21
add() (in module *Func.Math*), 20
and() (in module *Aggr.Meet*), 29
and() (in module *Func.Bool*), 20
append() (in module *Func.List*), 24
asin() (in module *Func.Math*), 21
asinh() (in module *Func.Math*), 21
assert() (in module *Func.Bool*), 20
atan() (in module *Func.Math*), 21
atan2() (in module *Func.Math*), 21
atanh() (in module *Func.Math*), 21

B

BetweennessCentrality() (in module *Algo*), 38
BFS() (in module *Algo*), 36
bit_and() (in module *Aggr.Meet*), 30
bit_and() (in module *Func.Bin*), 24
bit_not() (in module *Func.Bin*), 25
bit_or() (in module *Aggr.Meet*), 30
bit_or() (in module *Func.Bin*), 24
bit_xor() (in module *Aggr.Ord*), 30
bit_xor() (in module *Func.Bin*), 25
BreadthFirstSearch() (in module *Algo*), 35

C

ceil() (in module *Func.Math*), 20
chars() (in module *Func.String*), 23
choice() (in module *Aggr.Meet*), 29
choice_last() (in module *Aggr.Meet*), 30
chunks() (in module *Func.List*), 24
chunks_exact() (in module *Func.List*), 24
ClosenessCentrality() (in module *Algo*), 38
ClusteringCoefficients() (in module *Algo*), 37
coalesce() (in module *Aggr.Meet*), 30
collect() (in module *Aggr.Ord*), 30
CommunityDetectionLouvain() (in module *Algo*), 37
concat() (in module *Func.List*), 24
concat() (in module *Func.String*), 22
ConnectedComponents() (in module *Algo*), 34

cos() (in module *Func.Math*), 21
cosh() (in module *Func.Math*), 21
count() (in module *Aggr.Ord*), 30
count_unique() (in module *Aggr.Ord*), 30
CsvReader() (in module *Algo*), 33

D

decode_base64() (in module *Func.Bin*), 25
deg_to_rad() (in module *Func.Math*), 21
DegreeCentrality() (in module *Algo*), 37
DepthFirstSearch() (in module *Algo*), 36
DFS() (in module *Algo*), 36
difference() (in module *Func.List*), 24
div() (in module *Func.Math*), 20

E

encode_base64() (in module *Func.Bin*), 25
ends_with() (in module *Func.String*), 23
eq() (in module *Func.EqCmp*), 19
exp() (in module *Func.Math*), 20
exp2() (in module *Func.Math*), 21

F

first() (in module *Func.List*), 23
floor() (in module *Func.Math*), 20
format_timestamp() (in module *Func.Regex*), 28
from_substrings() (in module *Func.String*), 23

G

ge() (in module *Func.EqCmp*), 19
get() (in module *Func.List*), 23
group_count() (in module *Aggr.Ord*), 30
gt() (in module *Func.EqCmp*), 19

H

haversine() (in module *Func.Math*), 22
haversine_deg_input() (in module *Func.Math*), 22

I

intersection() (in module *Aggr.Meet*), 29
intersection() (in module *Func.List*), 24

is_bytes() (in module *Func.Typeing*), 26
 is_finite() (in module *Func.Typeing*), 26
 is_float() (in module *Func.Typeing*), 26
 is_in() (in module *Func.List*), 23
 is_infinite() (in module *Func.Typeing*), 26
 is_int() (in module *Func.Typeing*), 25
 is_list() (in module *Func.Typeing*), 26
 is_nan() (in module *Func.Typeing*), 26
 is_null() (in module *Func.Typeing*), 25
 is_num() (in module *Func.Typeing*), 26
 is_string() (in module *Func.Typeing*), 26
 is_uuid() (in module *Func.Typeing*), 26

J

JsonReader() (in module *Algo*), 34

K

KShortestPathYen() (in module *Algo*), 35

L

LabelPropagation() (in module *Algo*), 37
 last() (in module *Func.List*), 23
 latest_by() (in module *Aggr.Ord*), 30
 le() (in module *Func.EqCmp*), 19
 length() (in module *Func.Bin*), 24
 length() (in module *Func.List*), 23
 length() (in module *Func.String*), 22
 list() (in module *Func.List*), 23
 ln() (in module *Func.Math*), 21
 log10() (in module *Func.Math*), 21
 log2() (in module *Func.Math*), 21
 lowercase() (in module *Func.String*), 22
 lt() (in module *Func.EqCmp*), 19

M

max() (in module *Aggr.Meet*), 29
 max() (in module *Func.EqCmp*), 19
 maybe_get() (in module *Func.List*), 23
 mean() (in module *Aggr.Ord*), 31
 min() (in module *Aggr.Meet*), 29
 min() (in module *Func.EqCmp*), 19
 min_cost() (in module *Aggr.Meet*), 30
 MinimumSpanningForestKruskal() (in module *Algo*),
 34
 MinimumSpanningTreePrim() (in module *Algo*), 34
 minus() (in module *Func.Math*), 20
 mod() (in module *Func.Math*), 20
 mul() (in module *Func.Math*), 20

N

negate() (in module *Func.Bool*), 20
 neq() (in module *Func.EqCmp*), 19
 now() (in module *Func.Regex*), 28

O

or() (in module *Aggr.Meet*), 29
 or() (in module *Func.Bool*), 20

P

pack_bits() (in module *Func.Bin*), 25
 PageRank() (in module *Algo*), 37
 parse_timestamp() (in module *Func.Regex*), 28
 pow() (in module *Func.Math*), 20
 prepend() (in module *Func.List*), 24
 product() (in module *Aggr.Ord*), 31

R

rad_to_deg() (in module *Func.Math*), 21
 rand_bernoulli() (in module *Func.Rand*), 26
 rand_choose() (in module *Func.Rand*), 26
 rand_float() (in module *Func.Rand*), 26
 rand_int() (in module *Func.Rand*), 26
 rand_uuid_v1() (in module *Func.Rand*), 26
 rand_uuid_v4() (in module *Func.Rand*), 26
 RandomWalk() (in module *Algo*), 39
 regex_extract() (in module *Func.Regex*), 27
 regex_extract_first() (in module *Func.Regex*), 27
 regex_matches() (in module *Func.Regex*), 27
 regex_replace() (in module *Func.Regex*), 27
 regex_replace_all() (in module *Func.Regex*), 27
 ReorderSort() (in module *Algo*), 33
 reverse() (in module *Func.List*), 24
 round() (in module *Func.Math*), 20

S

SCC() (in module *Algo*), 34
 shortest() (in module *Aggr.Meet*), 30
 ShortestPathAStar() (in module *Algo*), 36
 ShortestPathDijkstra() (in module *Algo*), 35
 signum() (in module *Func.Math*), 20
 sin() (in module *Func.Math*), 21
 sinh() (in module *Func.Math*), 21
 slice() (in module *Func.List*), 24
 sorted() (in module *Func.List*), 24
 starts_with() (in module *Func.String*), 23
 std_dev() (in module *Aggr.Ord*), 31
 str_includes() (in module *Func.String*), 22
 StronglyConnectedComponent() (in module *Algo*), 34
 sub() (in module *Func.Math*), 20
 sum() (in module *Aggr.Ord*), 31

T

tan() (in module *Func.Math*), 21
 tanh() (in module *Func.Math*), 21
 to_float() (in module *Func.Typeing*), 25
 to_string() (in module *Func.Typeing*), 25
 to_uuid() (in module *Func.Typeing*), 25

`TopSort()` (*in module `Algo`*), [35](#)
`trim()` (*in module `Func.String`*), [22](#)
`trim_end()` (*in module `Func.String`*), [22](#)
`trim_start()` (*in module `Func.String`*), [22](#)

U

`unicode_normalize()` (*in module `Func.String`*), [23](#)
`union()` (*in module `Aggr.Meet`*), [29](#)
`union()` (*in module `Func.List`*), [24](#)
`unique()` (*in module `Aggr.Ord`*), [30](#)
`unpack_bits()` (*in module `Func.Bin`*), [25](#)
`uppercase()` (*in module `Func.String`*), [22](#)
`uuid_timestamp()` (*in module `Func.Typing`*), [25](#)

V

`variance()` (*in module `Aggr.Ord`*), [31](#)

W

`windows()` (*in module `Func.List`*), [24](#)